

3^η εργαστηριακή άσκηση προηγμένα θέματα αρχιτεκτονικής υπολογιστών

03117103 ΑΘΑΝΑΣΙΟΣ ΔΕΛΗΣ

A. Definitions(Wikipedia source)

Out of order superscalar processor: A superscalar processor is a [CPU](#) that implements a form of [parallelism](#) called [instruction-level parallelism](#) within a single processor. In contrast to a [scalar processor](#) that can execute at most one single instruction per clock cycle, a superscalar processor can execute more than one instruction during a clock cycle by simultaneously dispatching multiple instructions to different [execution units](#) on the processor.

In [computer engineering](#), out-of-order execution (or more formally dynamic execution) is a [paradigm](#) used in most high-performance [central processing units](#) to make use of [instruction cycles](#) that would otherwise be wasted. In this paradigm, a processor executes [instructions](#) in an order governed by the availability of input data and execution units,^[1] rather than by their original order in a program.^[2]

Issue/Dispatch width: issue width = the number of issue slots, 1 slot/instruction, concerns processors with the ability to execute Multiple instructions issued each cycle

Window size and ROB: An instruction window in [computer architecture](#) refers to the set of [instructions](#) which can [execute out-of-order](#) in a [speculative processor](#).



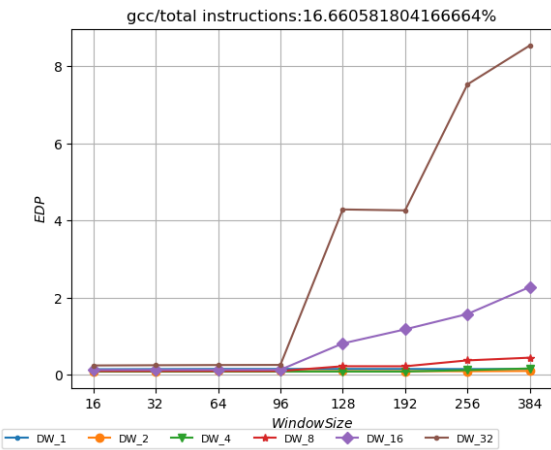
In particular, in a conventional design, the instruction window consists of all instructions which are in the [re-order buffer](#) (ROB).^[1] In such a processor, any instruction within the instruction window can be executed when its operands are ready. Out-of-order processors derive their name because this may occur out-of-order (if operands to a younger instruction are ready before those of an older instruction).

The instruction window has a finite size, and new instructions can enter the window (usually called dispatch or allocate) only when other instructions leave the window (usually called retire or commit). Instructions enter and leave the instruction window in program order, and an instruction can only leave the window when it is the oldest instruction in the window and it has been completed. Hence, the instruction window can be seen as a sliding window in which the instructions can become out-of-order. All execution within the window is speculative (i.e., side-effects are not applied outside the CPU) until it is committed in order to support asynchronous [exception handling](#) like [interrupts](#).


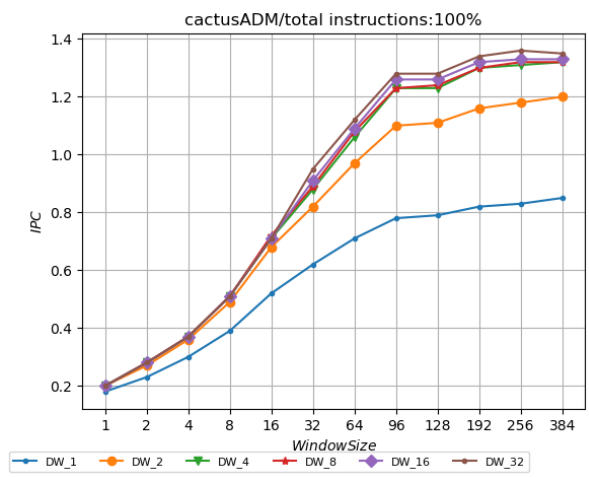
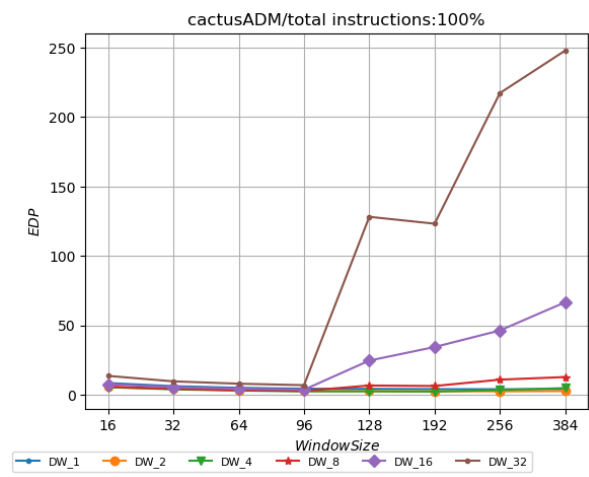
EDP: Metric runtime*enrgy

B.Plots


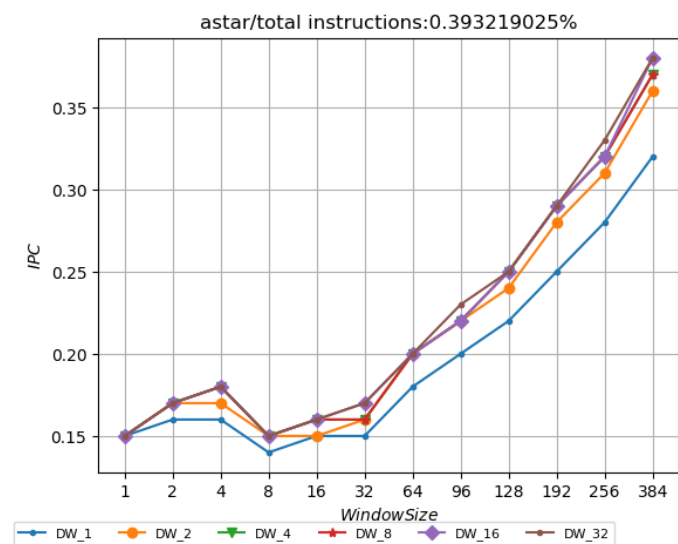
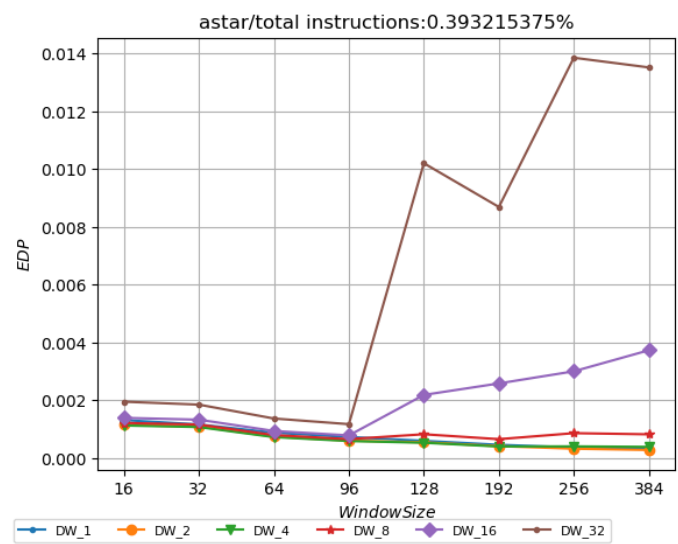
Gcc






CactusADM



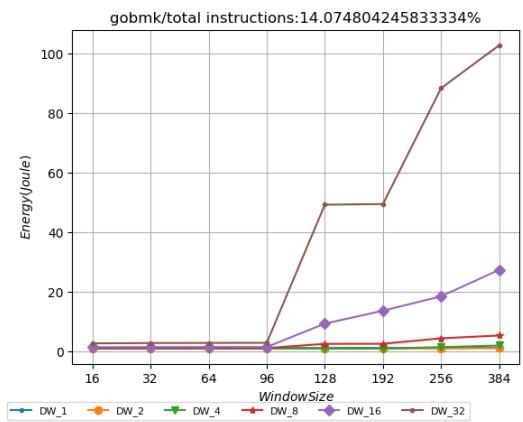
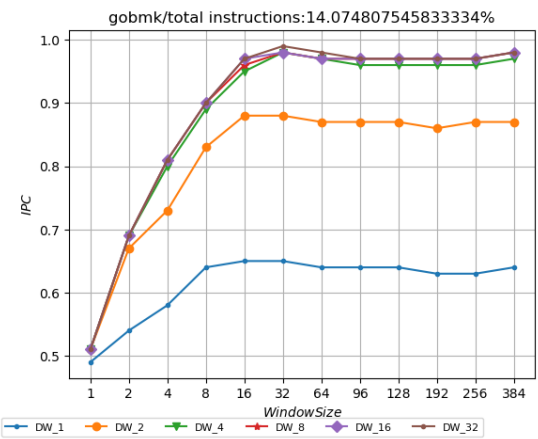

Astar




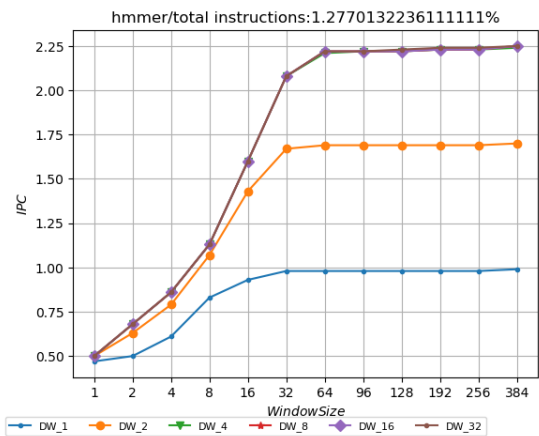
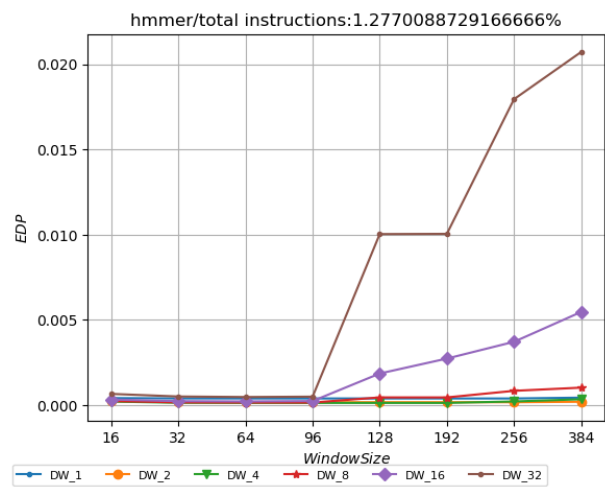
GemsFDTD




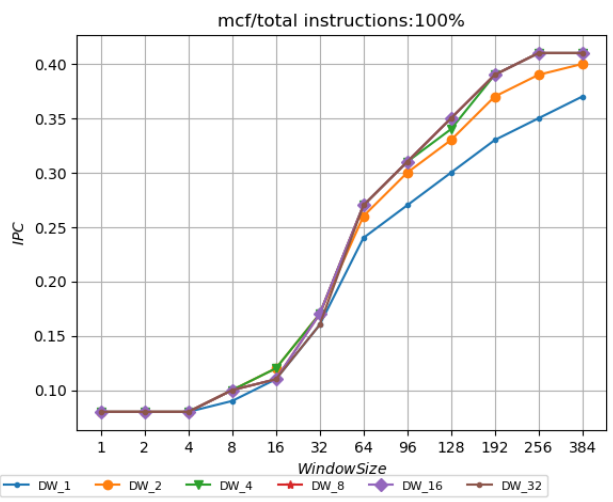
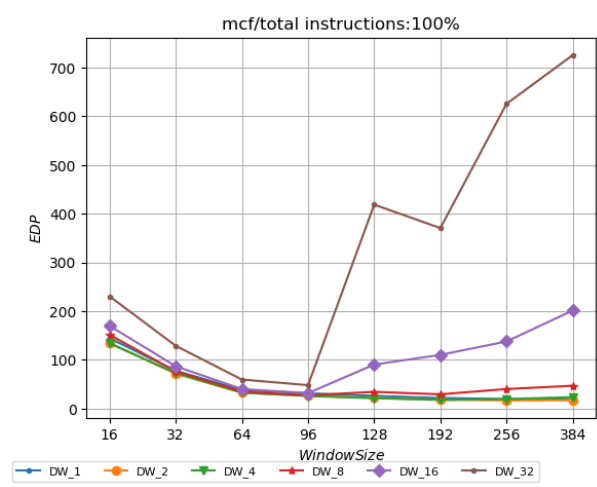
Gobmk



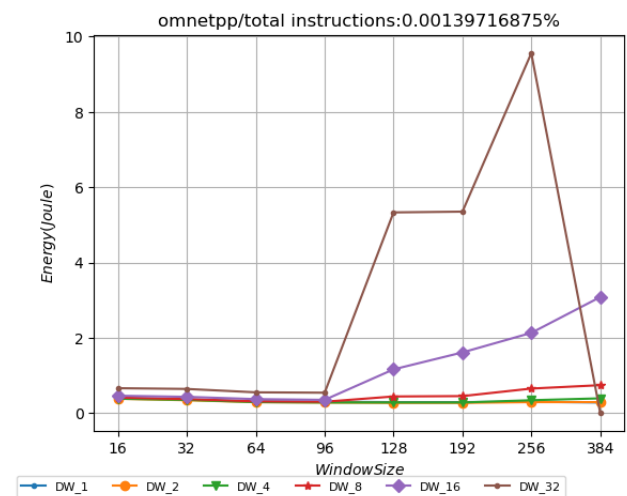
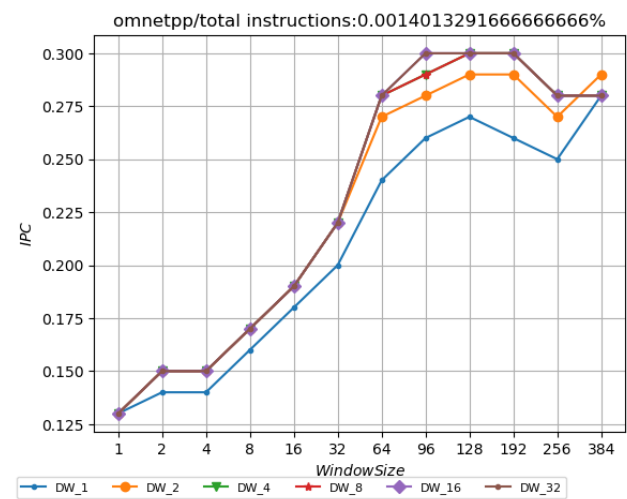

Hmmer





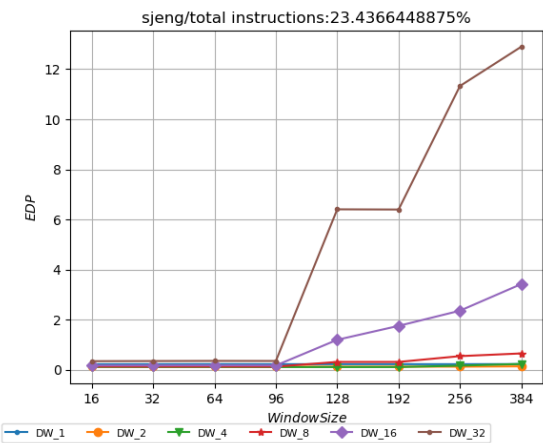
Mcf




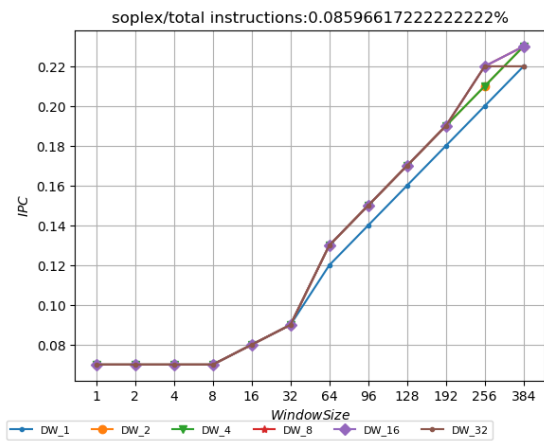

Omnetpp



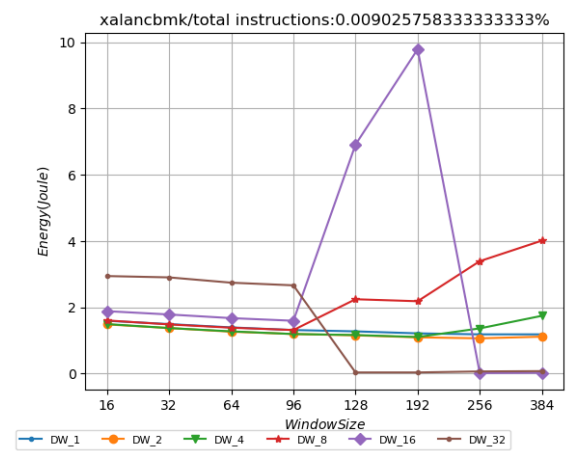
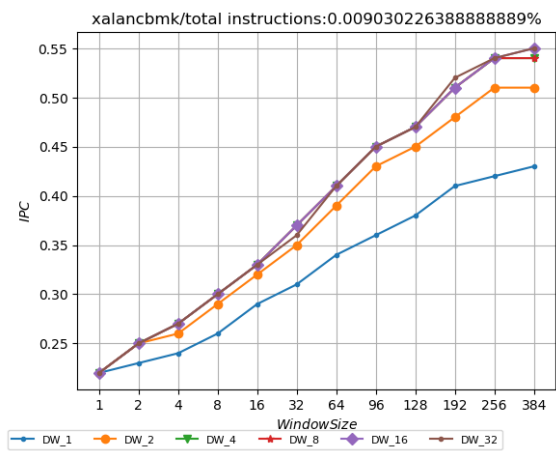

Sjeng



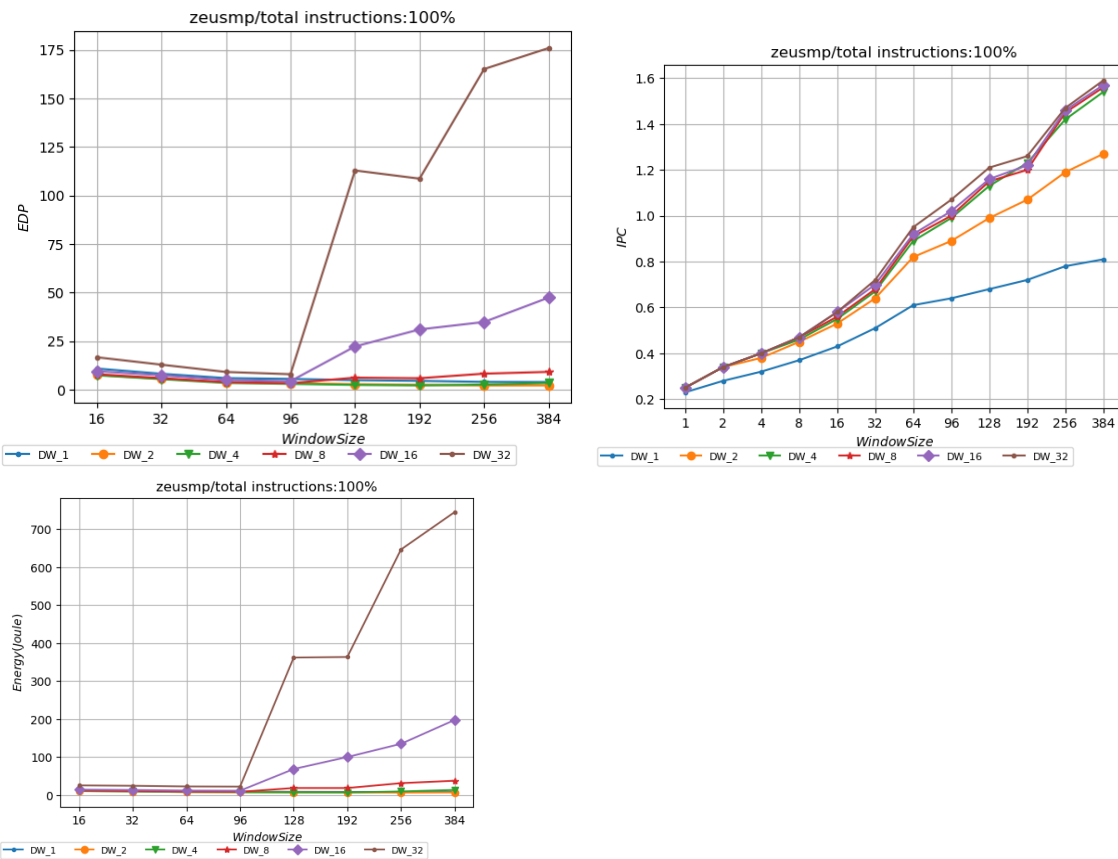
Soplex



Xalanbcbmk



Zeusmp



C. Questions-Conclusions

i) If the inequality window size < dispatch width is true, any further growth of dispatch width cannot alter the performance of our system. This is a case because in every clock cycle the processor will handle the workload that exists in the ROB, so anything more cannot be done even though we have further increase of the dispatch width. This would have allowed us to make less simulations in our study, except maybe some with window size < dispatch to be true in order to prove this observation. This approach was not followed though due to the circumstance that some of the benchmarks cannot be completed, due to internal failure, so our system resources could be used to simulate everything

ii) If we consider the function IPC with independent variable the window size we observe an increasing function, for different dispatch width values. So the increase of window size will give a better performance. In a narrow window of windows' values, a little different for its benchmark, the increase is exponential.

Another observation is that for the same window size value the increase of dispatch width results in better performance for our core (single core processors where simulated)

iii) As for the energy consumption of the processors we calculated a function of energy in joule with independent variable window size, for different dispatch width values, and again the function is increasing. We exclude some dispatch width values due to inefficient design of the McPAT program

The energy consumption is near zero (comparably to the rest of the function values really small) for window size values less than 96 (values of order 0 or 1, base of 10, no more than 100(J))

We calculated EDP, (in experimentation there was not a significant qualitative difference between EDP, ED^2P or other orders of ED^kP). In many, almost all, cases there was a decrease in EDP value till the value of window size 96 where exponential increase was observed.

IPC started to flatten in each benchmark, starting from values ranging from 8 to 64 with the mean value in between, no exceptions observed in the accepted benchmarks.

As a rule of thumb, we propose in a chip design window size between 16, 32, 64 to be implemented to achieve both performance and energy efficiency.

Dispatch width should be kept as low as possible to avoid extra unnecessary use of space, with no significant increase in performance, so we recommend between 2 to 8 (Possibly plots could be drawn though the lack of time could not allow that, this conclusion is driven from observation of some values and theory)

iv) Our processor: Intel core 2 cpu cpu 4400 2.00GHz but I could not find information for this processor

For SKYLAKE and SANDY BRIDGE I found that window size is 224 and 168 and dispatch width is 6 and 4 respectively. So for skylake we know that if ROB has 224 microoperations scheduler that has 8 ports can issue up to 8 instructions. Window size gives sufficient performance and we conclude that it is not wider due to energy restrictions

D. The code was written for energy information extraction and plotting of the various data, both .py and .sh files